

Oct 17, 20 20:20

hw5.c

Page 1/3

```

/* Homework Assignment #5
 * Simple CPU program to add two long vectors
 * and to calculate a 1D stencil
 * Author: Shawn Hinnebusch
 * Date: 10/19/2020
 *
 * Part 2 code is VECTORADD 1
 * Part 3-5 code is VECTORADD 0

 * To compile locally: nvcc -O3 -o hw5.exe hw5.cu -lm

 * To compile on the CRC:
 * crc-interactive.py -g -u 1 -t 1 -p gtx1080
 * nvcc -O3 -arch=sm_70 -o hw5.exe hw5.cu -lm

 * To run:
 * ./hw5.exe <vector size>
 * ./hw5.exe 1e8

 * Create PDF: a2ps hw5.cu âM-^HM-^RâM-^HM-^Rpro=color âM-^HM-^RâM-^HM-^Rcolumns
=2 âM-^HM-^RE âM-^HM-^Ro hw5.ps | ps2pdf hw5.ps

 */

#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <sys/resource.h>
#include "timer_nv.h"

#define VECTORADD 0 // else stencil

typedef float REAL;
typedef int INT;
#define RADIUS 3
#define BLOCK_SIZE 256

// ##### Stencil 1D Functions #####
__global__ void stencil_1d(const REAL *in, REAL *out, const INT n) {
    __shared__ REAL temp[BLOCK_SIZE + 2*RADIUS];
    INT gindex = threadIdx.x + blockIdx.x * blockDim.x;
    INT lindex = threadIdx.x + RADIUS;

    // Read input elements into shared memory
    __syncthreads();
    temp[lindex] = in[gindex];
    // Fills temp vector with previous value unless its the first block
    if (threadIdx.x < RADIUS && gindex > RADIUS) {
        temp[lindex - RADIUS] = in[gindex - RADIUS];
        // Check to not exceed the largest block size before filling in
        // the last 2 ghost cells
        if (gindex + BLOCK_SIZE < n) {
            temp[lindex + BLOCK_SIZE] = in[gindex + BLOCK_SIZE];
        }
    }
    __syncthreads();

    // Apply the stencil
    REAL result = 0.0;
    if (gindex >= RADIUS && gindex < (n - RADIUS)) {
        for (int offset = -RADIUS; offset <= RADIUS; offset++) {
            result += temp[lindex + offset];
        }
    }

    // Store the result

```

Oct 17, 20 20:20

hw5.c

Page 2/3

```

    out[gindex] = result;
}

void stencil_1d_cpu(const REAL *in, REAL *out, const INT n)
{
    for (int i = RADIUS; i < (n - RADIUS); i++) {
        for (int offset = -RADIUS; offset <= RADIUS; offset++) {
            out[i] += in[i + offset];
        }
    }
}

// ##### 2 Vector Functions #####

__global__ void vector_add_gpu (const INT n, const REAL *a, const REAL *b, REAL
*c)
{
    INT tid = blockIdx.x * blockDim.x + threadIdx.x;

    if (tid < n)
        c[tid] = a[tid] + b[tid];
}

void vector_add_cpu(const INT n, const REAL *a, const REAL *b, REAL *c)
{
    for (int i = 0; i < n; i++)
        c[i] = a[i] + b[i];
}

// ##### Main Function #####

int main(INT argc, char *argv[])
{
    if (argc < 2) {
        perror("Command-line usage: executableName <vector size>");
        exit(1);
    }

    int n = atof(argv[1]);

    printf("N: %d\n", n);

    // Initialize and alloc memory for arrays
    REAL *x, *y;
    cudaMallocManaged(&x, n * sizeof(*x));
    cudaMallocManaged(&y, n * sizeof(*y));

    #if VECTORADD
    REAL *z;
    cudaMallocManaged(&z, n * sizeof(*z));
    // Init vectors for addition
    for (int i = 0; i < n; i++) {
        x[i] = 3.5;
        y[i] = 1.5;
    }
    #else // stencil
    // Init vectors for stencil
    // note y is for the result should start at zero
    for (int i = 0; i < n; i++) {
        x[i] = 1.0;
        y[i] = 0.0;
    }
    #endif

    // CPU Run
    StartTimer();

```

Oct 17, 20 20:20

hw5.c

Page 3/3

```

    #if VECTORADD
    vector_add_cpu(n, x, y, z);
    #else // stencil
    stencil_ld_cpu(x, y, n);
    #endif

    double cpu_elapsedTime = GetTimer(); //elapsed time is in seconds

    // GPU Run
    cudaEvent_t timeStart, timeStop; //WARNING!!! use events only to time the de
vice
    cudaEventCreate(&timeStart);
    cudaEventCreate(&timeStop);
    float gpu_elapsedTime; // make sure it is of type float, precision is millis
econds (ms) !!!

    int nBlocks = (n + BLOCK_SIZE - 1) / BLOCK_SIZE; //round up if n is not a m
ultiple of blocksize

    cudaEventRecord(timeStart, 0); //don't worry for the 2nd argument zero, it i
s about cuda streams

    #if VECTORADD
    vector_add_gpu <<< nBlocks, BLOCK_SIZE >>> (n, x, y, z);
    #else // stencil
    stencil_ld <<< nBlocks, BLOCK_SIZE >>> (x, y, n);
    #endif

    cudaDeviceSynchronize();

    cudaEventRecord(timeStop, 0);
    cudaEventSynchronize(timeStop);

    cudaEventElapsedTime(&gpu_elapsedTime, timeStart, timeStop);

    printf("elapsed wall time (CPU) = %5.4f ms\n", cpu_elapsedTime*1000.);
    printf("elapsed wall time (GPU) = %5.4f ms\n\n", gpu_elapsedTime);

    cudaEventDestroy(timeStart);
    cudaEventDestroy(timeStop);

    //Used to print final results of CPU or GPU
    //for (int i = 0; i < n; i++) {
    //    printf("%f\n", y[i]);
    //}

    cudaFree(x);
    cudaFree(y);

    // Free additional vector for addition
    #if VECTORADD
    cudaFree(z);
    #endif

    return EXIT_SUCCESS;
}

```